

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df=pd.read_csv(r"C:\Users\USER\Downloads\Advertising.csv")
df
```

Out[2]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [3]: df.head()
```

Out[3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [4]: `df.tail()`

Out[4]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

In [5]: `df.shape`

Out[5]: (200, 4)

In [6]: `df.describe()`

Out[6]:

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV           200 non-null    float64
1    Radio        200 non-null    float64
2    Newspaper    200 non-null    float64
3    Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [8]: import seaborn as sns  
import matplotlib.pyplot as plt
```

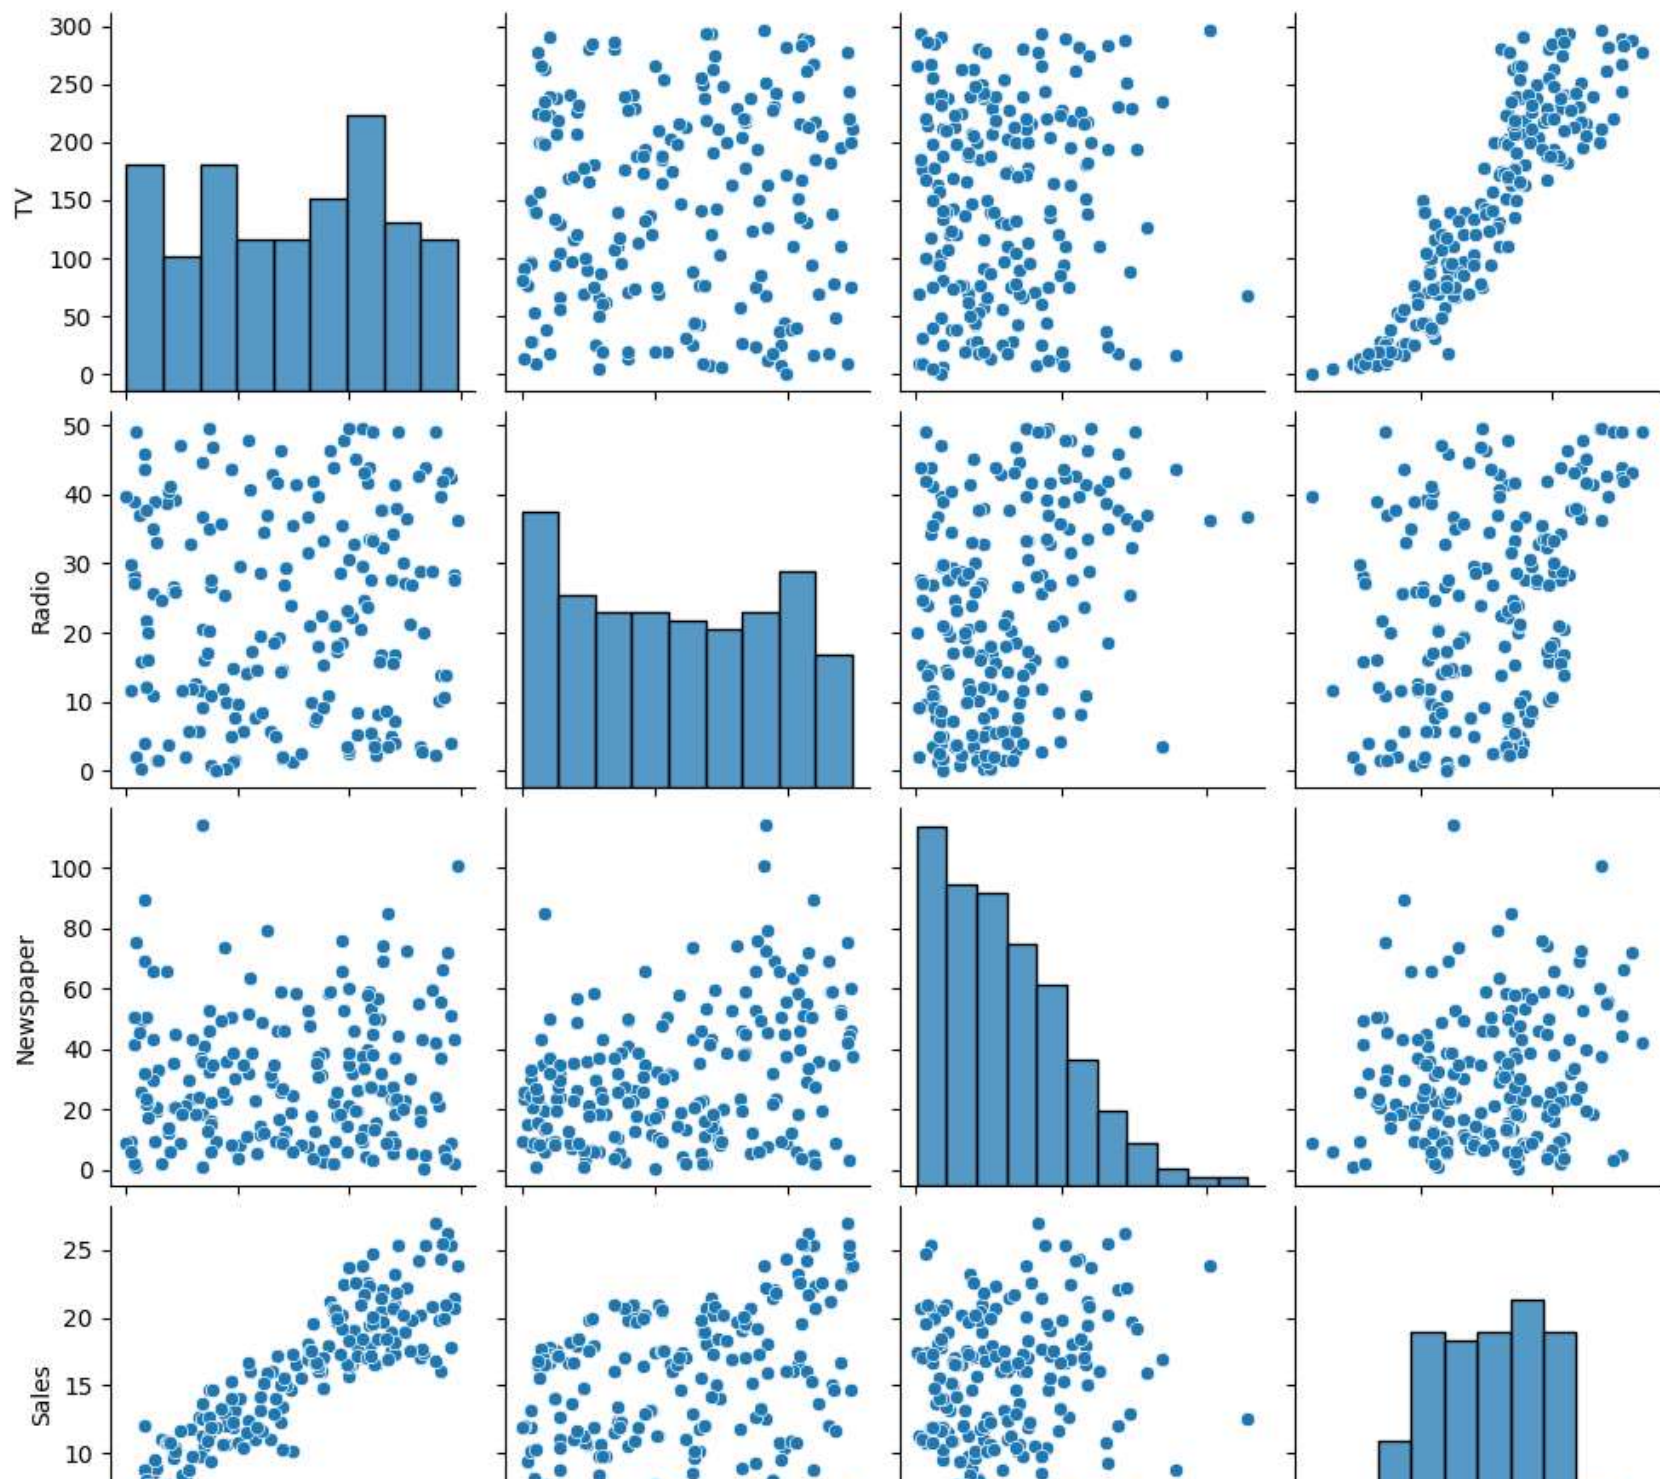
```
In [9]: df.isna().any()
```


```
Out[9]: TV          False  
Radio          False  
Newspaper      False  
Sales          False  
dtype: bool
```

```
In [10]: sns.pairplot(df)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x1dadbe57d50>
```









```
In [11]: sns.displot(df['Radio'])
```

```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x1dae5bc9110>
```



```
In [12]: sns.displot(df['TV'])
```

```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x1dae856e0d0>
```




```
In [13]: df.columns
```

```
Out[13]: Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```



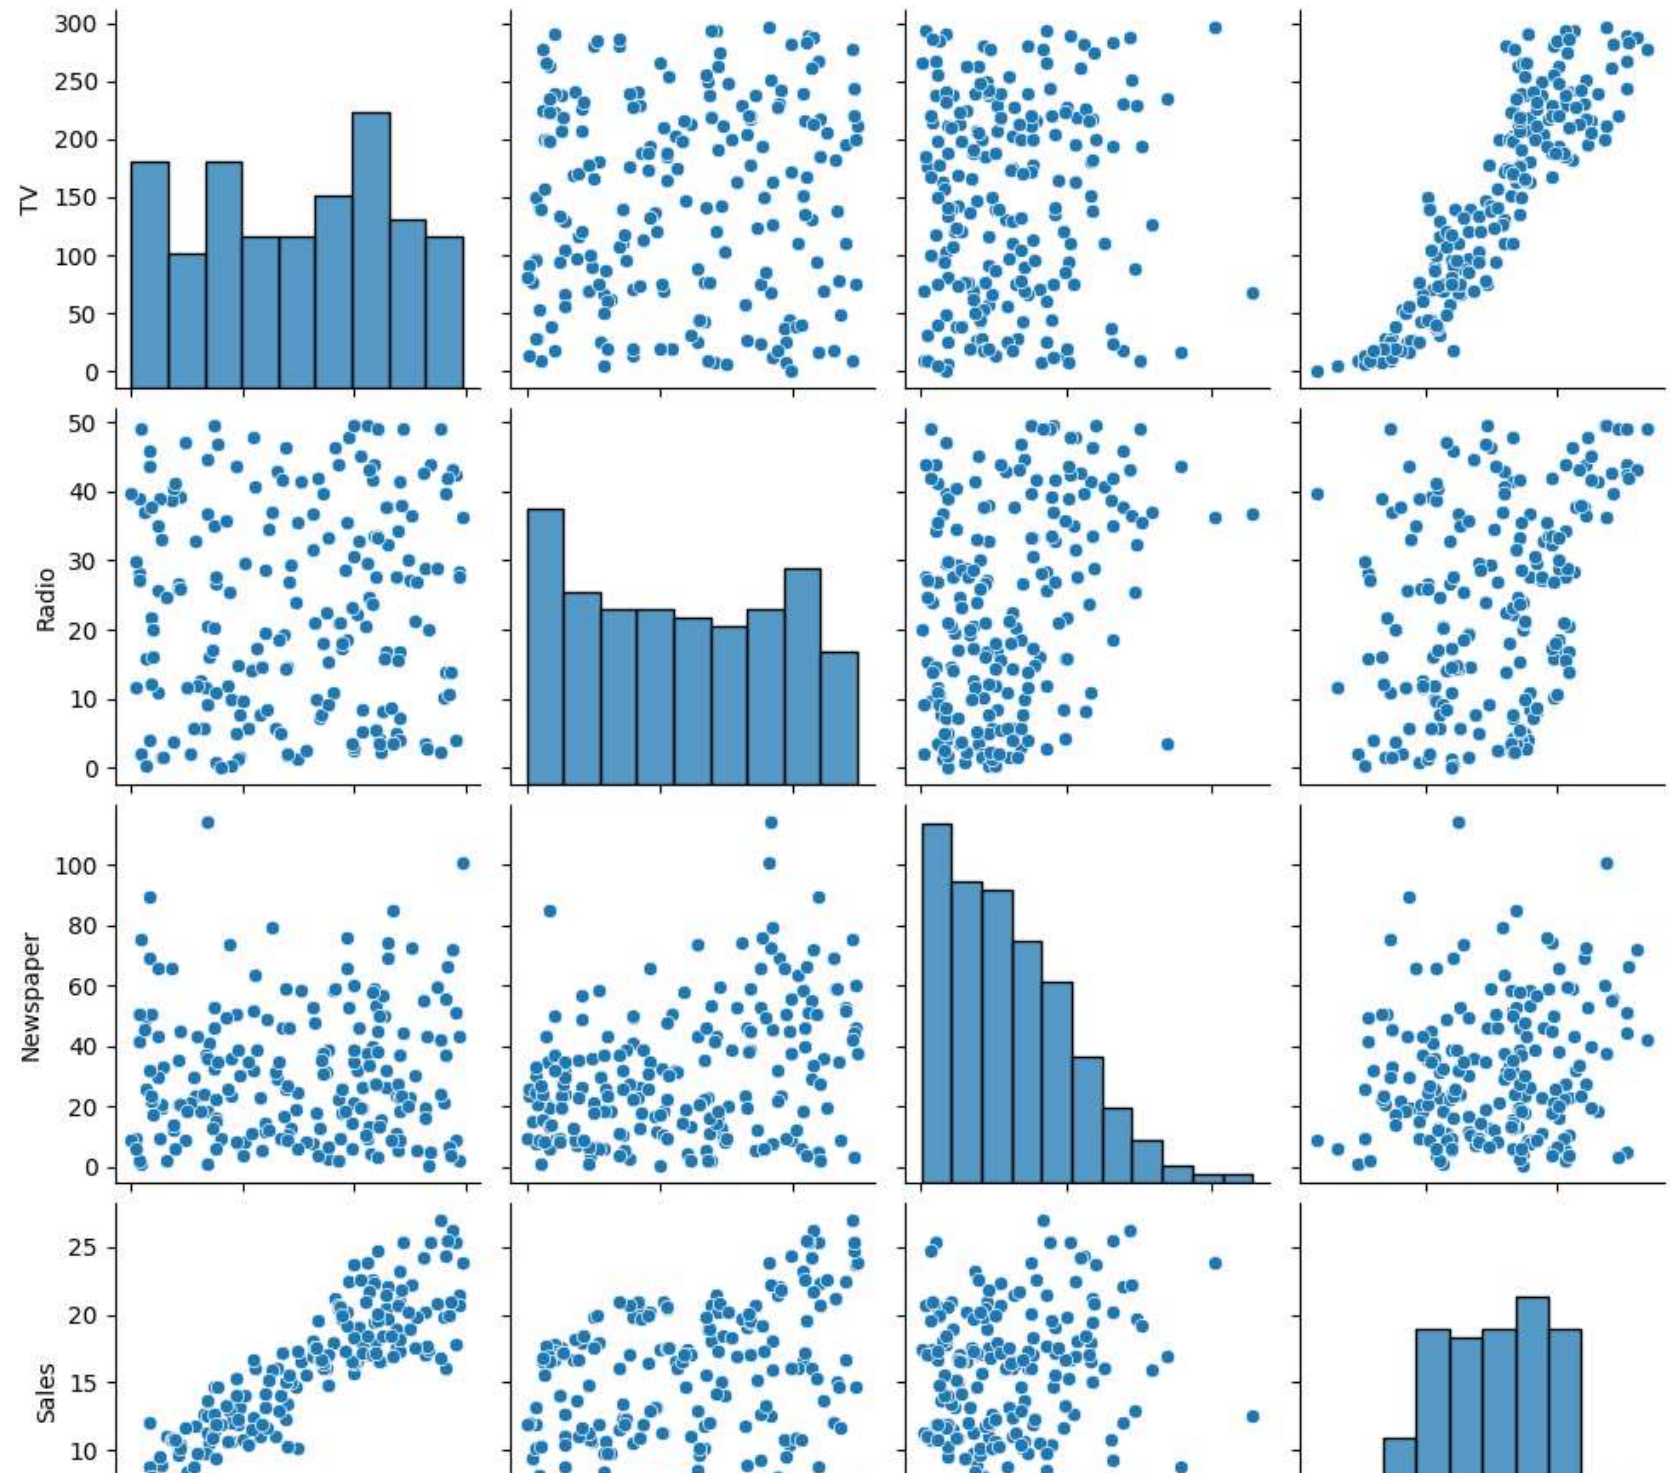
```
In [14]: sns.heatmap(df.corr())
```


```
Out[14]: <Axes: >
```



```
In [15]: df
sns.pairplot(df)
df.Sales = np.log(df.Sales)
```







```
In [17]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=101)
```

```
In [18]: from sklearn.preprocessing import StandardScaler
features=df.columns[0:2]
target=df.columns[-1]
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
Scaler=StandardScaler()
x_train=Scaler.fit_transform(x_train)
x_test=Scaler.transform(x_test)
```

The dimension of x\_train is (140, 2)

The dimension of x\_test is (60, 2)

```
In [19]: from sklearn.linear_model import LinearRegression
lm=LinearRegression()
lm.fit(x_train,y_train)
```


```
Out[19]: LinearRegression
LinearRegression()
```

```
In [20]: print(lm.intercept_)
```

2.649682499818669

```
In [21]: predictions=lm.predict(x_test)
plt.scatter(y_test,predictions)
```

Out[21]: <matplotlib.collections.PathCollection at 0x1daebcd8290>



```
In [22]: from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```


MAE: 0.13137245926116148

MSE: 0.0667845950252303

RMSE: 0.3624533890877025

```
In [23]: sns.displot((y_test-predictions),bins=50)
```

```
Out[23]: <seaborn.axisgrid.FacetGrid at 0x1daeaca3b10>
```




```
In [24]: print(lm.coef_)
```

```
[0.35159618 0.10807217]
```

```
In [25]: from sklearn.linear_model import Lasso,Ridge
```

```
In [26]: plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True)
```

Out[26]: <Axes: >



```
In [27]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```


Ridge model:

The train score for ridge model is 0.8500055285406656

The test score for ridge model is 0.6534250193956134



```
In [28]: df.drop(columns=["Radio", "Newspaper"], inplace=True)
sns.pairplot(df)
df.Sales=np.log(df.Sales)
```



```
In [29]: lr=LinearRegression()
lr.fit(x_train,y_train)
actual=y_test
train_score_lr=lr.score(x_train,y_train)
test_score_lr=lr.score(x_test,y_test)
print("\nLinear Regression model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression model:

The train score for lr model is 0.8534150366313791  
The test score for lr model is 0.6654095499889279

```
In [30]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.8500055285406656

The test score for ridge model is 0.6534250193956134

```
In [31]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [32]: print("\nLasso model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso model:

The train score for ls model is 0.0

The test score for ls model is -0.0042092253233847465

```
In [33]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```


Out[33]: <Axes: >



```
In [34]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.8534149297347071  
0.6654001015086553

```
In [36]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*', markersize=5, color='red',label=r'Ridge;  $\alpha=10$ ', zorder=1)
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha=grid$ ')
plt.plot(features, lr.coef_, alpha=0.4, linestyle='none', marker='o', markersize=7, color='green',label='Linear Regression')
plt.xticks(rotation=90)
plt.legend()
plt.title('Comparison plot of Ridge,Lasso and LinearRegression model')
plt.show()
```



```
In [38]: #Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(x_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(x_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(x_test, y_test)))
```

The train score for ridge model is 0.8534146479788303  
The train score for ridge model is 0.6653038057268583

```
In [48]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x_train,y_train)
print(regr.coef_)
print(regr.intercept_)
```

[0. 0.]  
2.649682499818669

```
In [52]: y_pred_elastic=regr.predict(x_train)
```

```
In [55]: mean_squared_error=np.mean((y_pred_elastic- y_train)**2)
print("mean _squared_error",mean_squared_error)
```

mean \_squared\_error 0.16840246163748074

```
In [57]: features = df.columns[0:2]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X\_train is (140, 2)  
The dimension of X\_test is (60, 2)

```
In [58]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```


Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0

```
In [64]: lt.figure(figsize = (10, 10))
lt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha = 10$ ',zorder=7)
plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;  $\alpha = 100$ ')
lt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
lt.xticks(rotation = 90)
lt.legend()
lt.show()
```





In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: